

User tutorial

A detailed, hands-on guide to `matirf_python_plus`: the two working modes, every control of the GUI, the file types in and out, running a reconstruction, and reading the results against a ground truth. For installation see the [README](#); for the theory and per-algorithm guidance see [docs/algorithms/](#).

Screenshots use the default **dark theme** (`settings.toml`: `dark_style = true`, `app_style = "fusion"`). Yours match your own settings.

1. The idea, and the two modes

The tool reconstructs an unknown object $\mathbf{f}$ from a measurement $\mathbf{g}$ and a known forward operator $\mathbf{H}$, by minimising $(1 - \lambda) \cdot D(\mathbf{H}\mathbf{f}, \mathbf{g}) + \lambda \cdot R(\mathbf{f})$ under $\mathbf{f} \geq \mathbf{0}$. Two inverse problems ship with it — **matirf** (3D MA-TIRF depth reconstruction) and **deconv** (2D deconvolution). Every command below works for both: swap `matirf` for `deconv`.

There are **two modes**, chosen by a toggle at the top of the control window:

- **Synthetic** — you provide a known object (a *ground truth* `f_true`); the tool simulates the measurement $\mathbf{g} = \mathbf{H} \cdot \mathbf{f}_{\text{true}}$ (optionally adding noise), reconstructs, and **compares the result to the truth** (metrics, §6). Use this to test and tune algorithms on data whose exact answer you know.
- **Real** — you provide a real measurement $\mathbf{g}$; the tool reconstructs $\mathbf{f}$. There is no truth, so no error metrics — only the data-fit diagnostic and the visual result.

2. The control window

`matirf` gui




Figure 1: The `matirf` control window (synthetic mode)

1. **Input Files.** The measurement (or, in synthetic mode, the object) and the microscope parameters. See the file types in §3. **See preprocessed file** shows g exactly as the reconstruction will see it. **Normalization** fixes the unit g (and therefore the noise level and the data term) is expressed in — peak ($g \leftarrow g/\max g$) is the default.
2. **Mode toggle.** Switch between *real measurement* and *simulate from a synthetic truth* (§1). The Input Files labels and the availability of the noise-injection section follow it.
3. **Operator Parameters.** The depth grid H reconstructs on: number of planes n_z , and the depth range z_0 – z_N (nm). **Normalize Operator** stays *off* for MA-TIRF (it would divide H by its largest singular value; the physical operator is wanted). See [00_foundation.md](#).
4. **Add noise to measurement** (*synthetic only*). Inject Gaussian read noise and/or Poisson shot noise with a fixed **seed** (same seed $\rightarrow$ same noise, for reproducible comparisons). The formula shown is exactly what is applied.
5. **Noise model (data fidelity).** The likelihood D the reconstruction assumes — Gaussian (L2), Poisson, or Poisson-Gaussian. **Noise level (a, b):** *estimated* reads it from g ; *oracle* (synthetic only) uses the exact level you injected. The formula and the current (a , b) are shown live.
6. **Select Algorithm.** Pick the solver (ADAM, PPXA, ADMM, PnP, ADMM-PnP, MCMC); the parameter panel adapts to it. **reset parameters** restores its defaults. **Estimate** buttons compute data-driven values (here the anisotropy ratio δ , and the ridge start weight λ_{rr}). What each parameter does — and its useful range — is in [docs/algorithms/](#).
7. **Load any config / Save config.** Read or write the whole setup as a `.toml` (§3).
8. **Open reconstruction.** Re-open a previously saved result in the display window.
9. **Run.** Start the reconstruction; the display window (§5) opens.

3. File types — input and output

Inputs

file	what it is
.tif/.TIF	the 3D data — a synthetic object <code>f_true</code> (synthetic mode) or the measurement stack <code>g</code> (real mode). For deconv it is a <code>*.png</code> image.
*.json	the microscope / measurement parameters (angles, wavelength, refractive indices, NA...) that define the operator H . Modify .json file edits it in place.
config.toml	the whole run — inputs, operator, noise, algorithm and its parameters. This is what <i>Save config</i> writes and <i>cli</i> reads.

A `config.toml` looks like this:

```
algorithm = "ADAM"
```

[input-paths]

```
mode = "synthetic-data"           # or "real-data"
tif = ".../cell_fibres_vesicles.TIF"
json = ".../measurement_parameters.json"
normalization = "peak"
```

[oper-params]

```
nz = 50
z0 = 0.0
zN = 300.0
normalize = false
```

[add-noise]

```
# synthetic only; empty {} for real data
```

```

gaussian_noise = true
sigma = 0.02
seed = 1

[noise-model]
parameters = "oracle"           # or "estimated"

[algo-params]
reg = "L1 norm of the gradient" # the regulariser R
lambda_reg = 0.1                # the regularisation share  $\lambda$ 
max_iter = 3000
K = 10
EPS = 1e-6

```

Outputs — `Save reconstruction` (or `cli -o`) writes a folder:

file	what it is
config.toml	the exact run (re-run it to reproduce)
f.tif (f.png for deconv)	the reconstruction
f_true.tif	the truth (synthetic mode)
metrics.json	the quality metrics (synthetic mode, §6)
messages.txt	the run log

4. Real vs synthetic mode

Toggling to **real** changes the Input Files: you now point to your **measurement stack** (not a truth) and its parameters, and the noise-injection section no longer applies (a real measurement already carries its noise).

The rest is identical. The one consequence downstream: with no truth there is **no metrics table** — only the data-fit diagnostic (`chi2_ratio`, which needs no truth) and the visual result.

5. Running, and the display window

Pressing **Run** opens the display window. On the left, the **Config** (the exact setup) and the **Messages** log (iterations, runtime, and the *realism check* verdict). On the right, the figures — with a **view 1 / view 2** toggle.

View 1 — the depth map

1. **Config** — the exact setup this result came from (the `config.toml`).
2. **Messages** — the iteration log, the execution time, and the realism check (here a mild *depth-shift* warning).
3. **Figures (view 1)** — the **Depths map** (each lateral pixel coloured by the depth, in nm, of its brightest voxel — the MA-TIRF answer, “how far from the glass”), and the **yz / zx profiles** (cross-sections through the volume). The matplotlib toolbars pan/zoom/save each.
4. **view 1 / view 2 toggle** — switch between the depth-map view (view 1, here) and the intensity + histogram view (view 2, below).
5. **Save view as PNG** — export this figure at a fixed size (used for the report figures).
6. **Save reconstruction** — write the output folder (§3).
7. **Go To Truth** — switch to the ground-truth comparison (§6). **Synthetic mode only** — in real mode there is no truth, so this button is absent.




Figure 2: The control window in real mode




Figure 3: Display window, view 1: depth map and profiles




Figure 4: Display window, view 2: intensity slices and 3D histogram

View 2 — intensity and histogram

View 2 shows the raw **intensity** of the reconstruction: a slice viewer (scroll through the nz planes), the pixel read-out, and the **histogram of the whole volume**. Note the values sit around $0-0.02$ — on MA-TIRF f peaks near 0.02 , not 1 (this is why parameters “in units of f ” must be read on that scale; see [00_foundation.md](#)).

6. Metrics and ground-truth analysis

In synthetic mode, **Go To Truth** shows the **Synthetic Ground Truth Analysis** — every metric comparing the reconstruction to the known f_{true} :

- **Scale_alpha** — the positive factor α that best matches f to the truth. MA-TIRF fixes f only up to such a factor (here $\alpha \approx 37 \approx s1$), so every error metric is computed **after** aligning by α .
- **NMSE** — normalised mean-squared error (0 perfect, 1 = no better than zero). The workhorse.
- **PSNR / SSIM** — the familiar image-processing measures (PSNR in dB; SSIM structural).
- **Depth Error (nm)** — the axial-localisation error, **the** MA-TIRF question.
- **Stack Recovery** — the fraction of ≥ 30 nm-separated structure pairs the reconstruction also resolves (axial super-resolution).
- MSE/MAE, angular distance, cosine similarity, correlation, Sinkhorn, FSC — additional reference measures.
- **Visualize difference** opens $f_{\text{true}} - \alpha \cdot f$ to see *where* the reconstruction errs.

(The benchmark in benchmarks/ records a curated subset of these — NMSE, PSNR, Depth Error, Stack Recovery, plus χ^2_{ratio} which needs no truth; see [docs/benchmark/](#).)




Figure 5: The metrics table vs the ground truth

7. Making synthetic test data

To compare algorithms on data whose exact answer is known, design a ground truth:

`matirf synth`

Set the grid (voxels, planes, depth range) and the structures (ellipsoids for vesicles/adhesions, filaments, membranes, double layers), then:

1. **Generate / preview** — build and show the truth.
2. **Save truth...** — write it as a `.tif` plus a `.truth.json` that reproduces it bit-for-bit.
3. **Use as MA-TIRF truth** — load it straight into the `matirf` config as the synthetic object (then just press Run in the control window).
4. **Load scene... / 5. Save scene...** — store the whole design as a reusable `.toml`.

8. Headless / batch (CLI)

Run a reconstruction with no interface, from a saved config:

`matirf cli -c my_config.toml -o results/run-01`

It prints progress and writes the output folder (§3). This is how to script many runs; the campaign in `benchmarks/` does exactly this at scale (see [docs/benchmark/](#)).

9. Settings

```

settings show           # list every setting and its value
settings set dark_style false # e.g. switch to the light theme

```




Figure 6: The synthetic ground-truth generator

```
settings reset           # back to defaults
settings gui             # edit them in a window
```

Useful ones: `dark_style` / `app_style` (appearance), `font_size_*`, `device` (cpu / cuda / mps), `dtype`, and `live_preview` (refresh the figures during the run).

10. Troubleshooting

Symptom → likely cause → what to do. (For *what a parameter does* and its useful range, see [docs/algorithms/](#) — this table is only about things going wrong.)

symptom	likely cause	what to do
<i>“Realism check REJECTED: depth profile far from...”</i>	the reconstruction is implausible (emptied, collapsed or depth-shifted) — usually too much or too little prior	adjust λ / κ ; match the prior to the structure (sparsity for point-like objects, a denoiser for continuous ones)
warning <i>“depth profile shifted by ...”</i>	noise with too weak a prior: the depth drifts	raise the regularisation — on MA-TIRF the prior, not the data, decides the depth
reconstruction is empty / flat	over-regularised ($\lambda \rightarrow 1$, or a high ADMM κ that empties the image)	lower λ / κ
out-of-memory / timeout on real data	large real volume with a costly prior (SHV) or the PnP family	avoid SHV on full real volumes; use a lighter prior; reduce the resolution
the run fails at start	file not found, invalid <code>.json</code> , or mismatched dimensions (angles vs stacks)	the error message says which; check the paths and the <code>.json</code>

symptom	likely cause	what to do
reconstruction is very slow	device = cpu, a large volume, many iterations	use a GPU (settings set device cuda/mps); the EPS criterion stops once the loss plateaus
result at a different scale than the truth	expected — MA-TIRF fixes f only up to a positive factor α	nothing; metrics align by α first (see Scale_alpha, §6)

11. Housekeeping & next steps

- `matirf reset` deletes the cached `config.toml` (back to defaults).
- `list` shows the available inverse problems; `help` shows every command.

Next: the decision guide and per-algorithm notes in [docs/algorithms/](#), and the [developer guide](#) (adding an inverse problem, a solver, a metric).